

# Towards enhanced performance and energy savings for new Cloud services

DJOB MVONDO

UNIVERSITY OF RENNES 1-IRISA-INRIA, FRANCE

21-09-2022

# CONTEXT : FUNCTION AS A SERVICE CLOUD MODEL

Serverless cloud model is gaining a lot of traction.

~ 22 Billion \$ estimated by 2025<sup>1</sup>

Amazon Lambda

Google Functions


Azure Functions

...

[1] <https://www.alliedmarketresearch.com/serverless-architecture-market>

# CONTEXT: FUNCTION AS A SERVICE CLOUD MODEL

Developers send the **code** and configures the **events/trigger**



Billed based on execution time and memory used.




Focus on your code and leave the rest to the provider

## CONTEXT : FUNCTION AS A SERVICE CLOUD MODEL

The cloud scheduler will direct the request to a server to launch the isolation units.


On each server, the OS scheduler must ensures **fair sharing** of CPU time for every isolation unit



# PROBLEM: **IDLE** ISOLATION UNITS

However, some isolation units may be idle

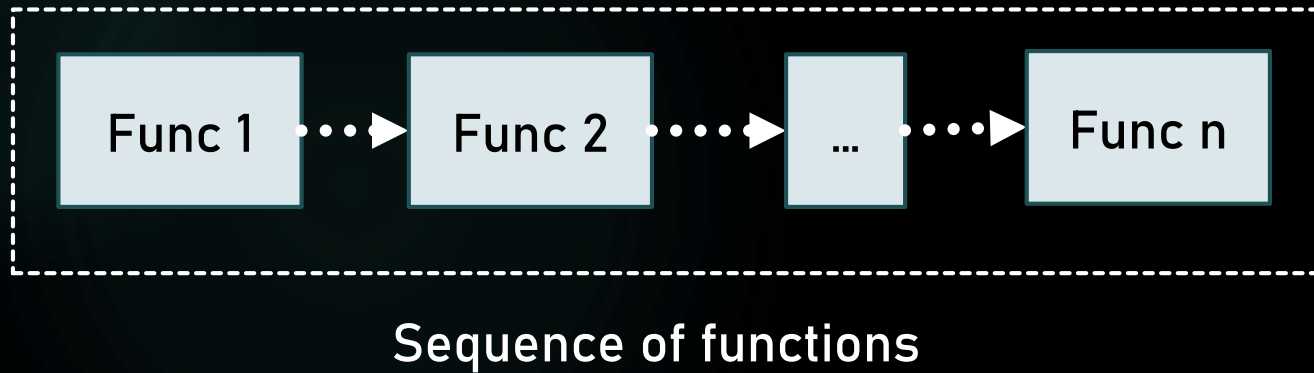
- **Keep alive policy** to reduce functions' start-up time.



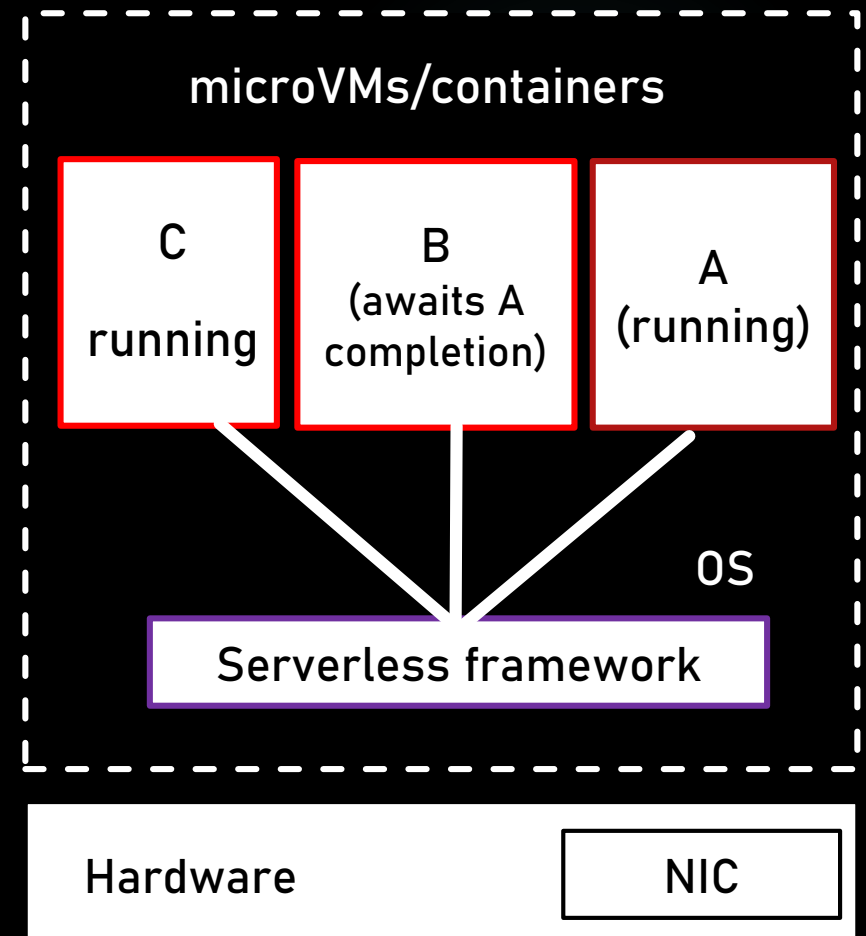
# PROBLEM : **IDLE** ISOLATION UNITS

However, some isolation units may be **idle**

- **Functions awaiting inputs** from other functions



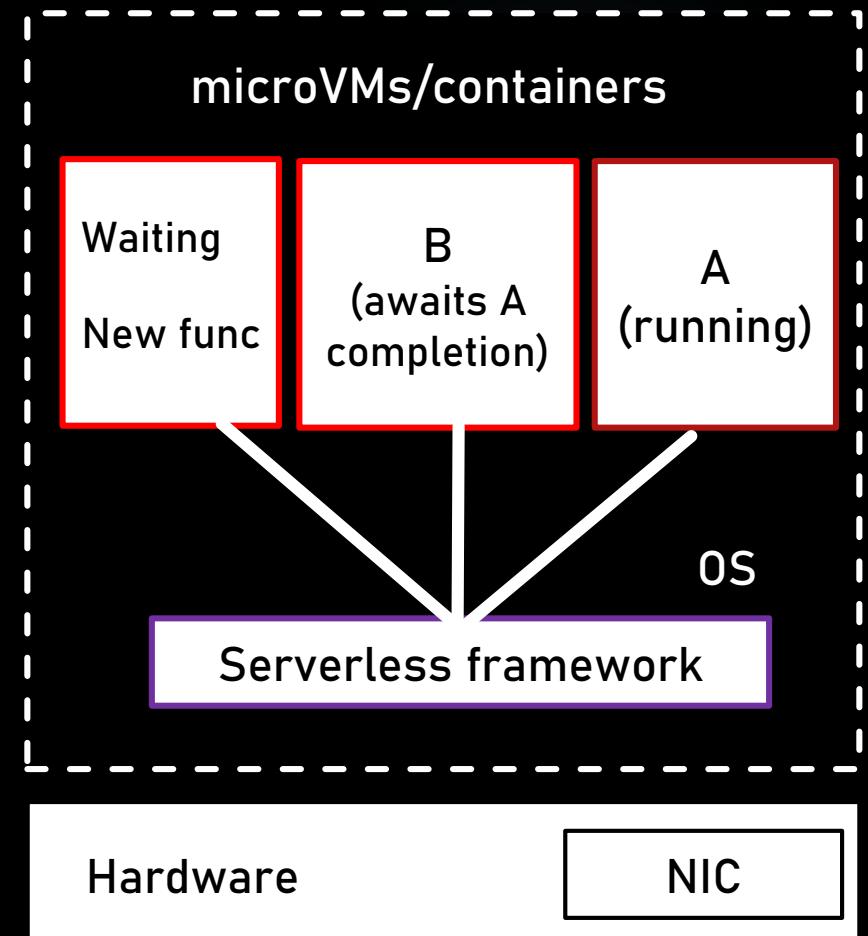
All isolation units are triggered at the same time to reduce cold latencies



# PROBLEM: **IDLE** ISOLATION UNITS

The isolation units idleness raises two main issues


1. **Wasted CPU time**
2. **Incorrect accounting misleading CPU frequency variations**



## PROBLEM 1: WASTED CPU TIME ON IDLE ISOLATION UNITS

For a sequence of 3 functions,  $\text{Func}\{1,2,3\}$ .

$\text{Func}\{2,3\}$  isolation units are initialized but await func1 completion



**Figure 1.** Illustration of micro-VMs idle times. Micro-VMs *b* and *c* running *Func2* and *Func3* respectively, are scheduled even though they await *Func1* output which has not finished running. This results in wasted CPU time.



## PROBLEM 1: **WASTED CPU TIME** ON IDLE ISOLATION UNITS

We analyzed the wasted CPU time on idle isolation units.

In-lab setup and ec2 a1.metal with Firecracker<sup>2</sup>

Triggering up to 50 pipelines image processing functions

Inputs and outputs images stored in AWS S3

We compute isolation units idle CPU usage

[2] Alexandru Agache et al. Firecracker: Lightweight Virtualization for Serverless Applications NSDI'20

# PROBLEM: **WASTED CPU TIME** ON IDLE ISOLATION UNITS

Avg Pipeline execution time

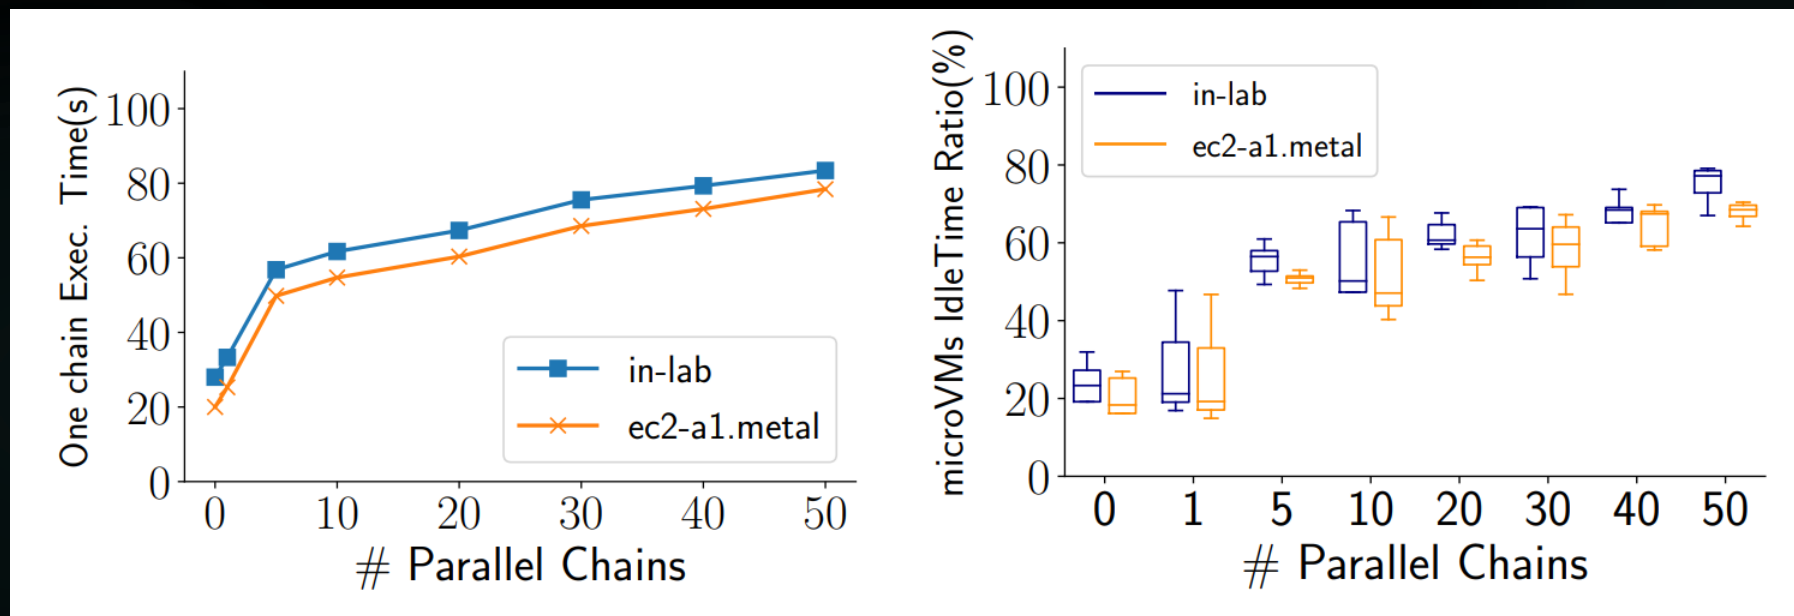
28.3s to 83.41s --- inlab

20s to 78.52s --- AWS a1.metal

Idle time ratio

20.18% to 75.31% --- inlab

16.25% to 69% --- AWS a1.metal



CPU time is **wasted**, smarter use could **improve** overall execution time

# Possible ideas: Scheduling semantic gap

Well known problem in the context of virtualization.

The host scheduler should :

Understand when an isolation unit  
is idle

- How do you detect ?
- What's the penalty of a false positive ?
- Intrusiveness ?

Understand the events that will  
affect idle isolation units

- Monitor events
- Dynamically update scheduling policy ?

# Possible ideas : Scheduling semantic gap

## Approaches worth exploring

Understand when an isolation unit  
is idle

How do you detect ?

What's the penalty of a  
false positive ?

Intrusiveness ?



Trusted source

Collaborative

Learning

## Possible ideas: Scheduling semantic gap

Well known problem in the context of virtualization.

The host scheduler should :

Understand the events that will  
affect idle isolation units

Monitor events

Dynamically update  
scheduling policy ?

# Problem : Scheduling bad performance

Dynamic scheduler behavior...

Not an easy task

Fixed policies (extensible to some extent but remain rigid)

Patching, scheduler class with a light interface

Lack of visibility on users' thread behavior

*Justinien Bouron et al. Thee Battle of the Schedulers: FreeBSD ULE vs. Linux CFS. ATC 2018*

*Redha Gouicem et al. Fewer Cores, More Hertz: Leveraging High-Frequency Cores in the OS Scheduler for Improved Application Performance. ATC 2020*

*Weiwei Jia et al. vSMT-IO: Improving I/O Performance and Efficiency on SMT Processors in Virtualized Clouds. ATC 2020*

*Bao Bui et al. When eXtended Para-Virtualization (XPV) meets NUMA. Eurosys 2019*

*Jean Pierre Lozi et al. . The Linux scheduler: a decade of wasted cores. Eurosys 2016*

# Our idea: Towards user-programmable schedulers

Users **express** the scheduler behavior at **runtime**

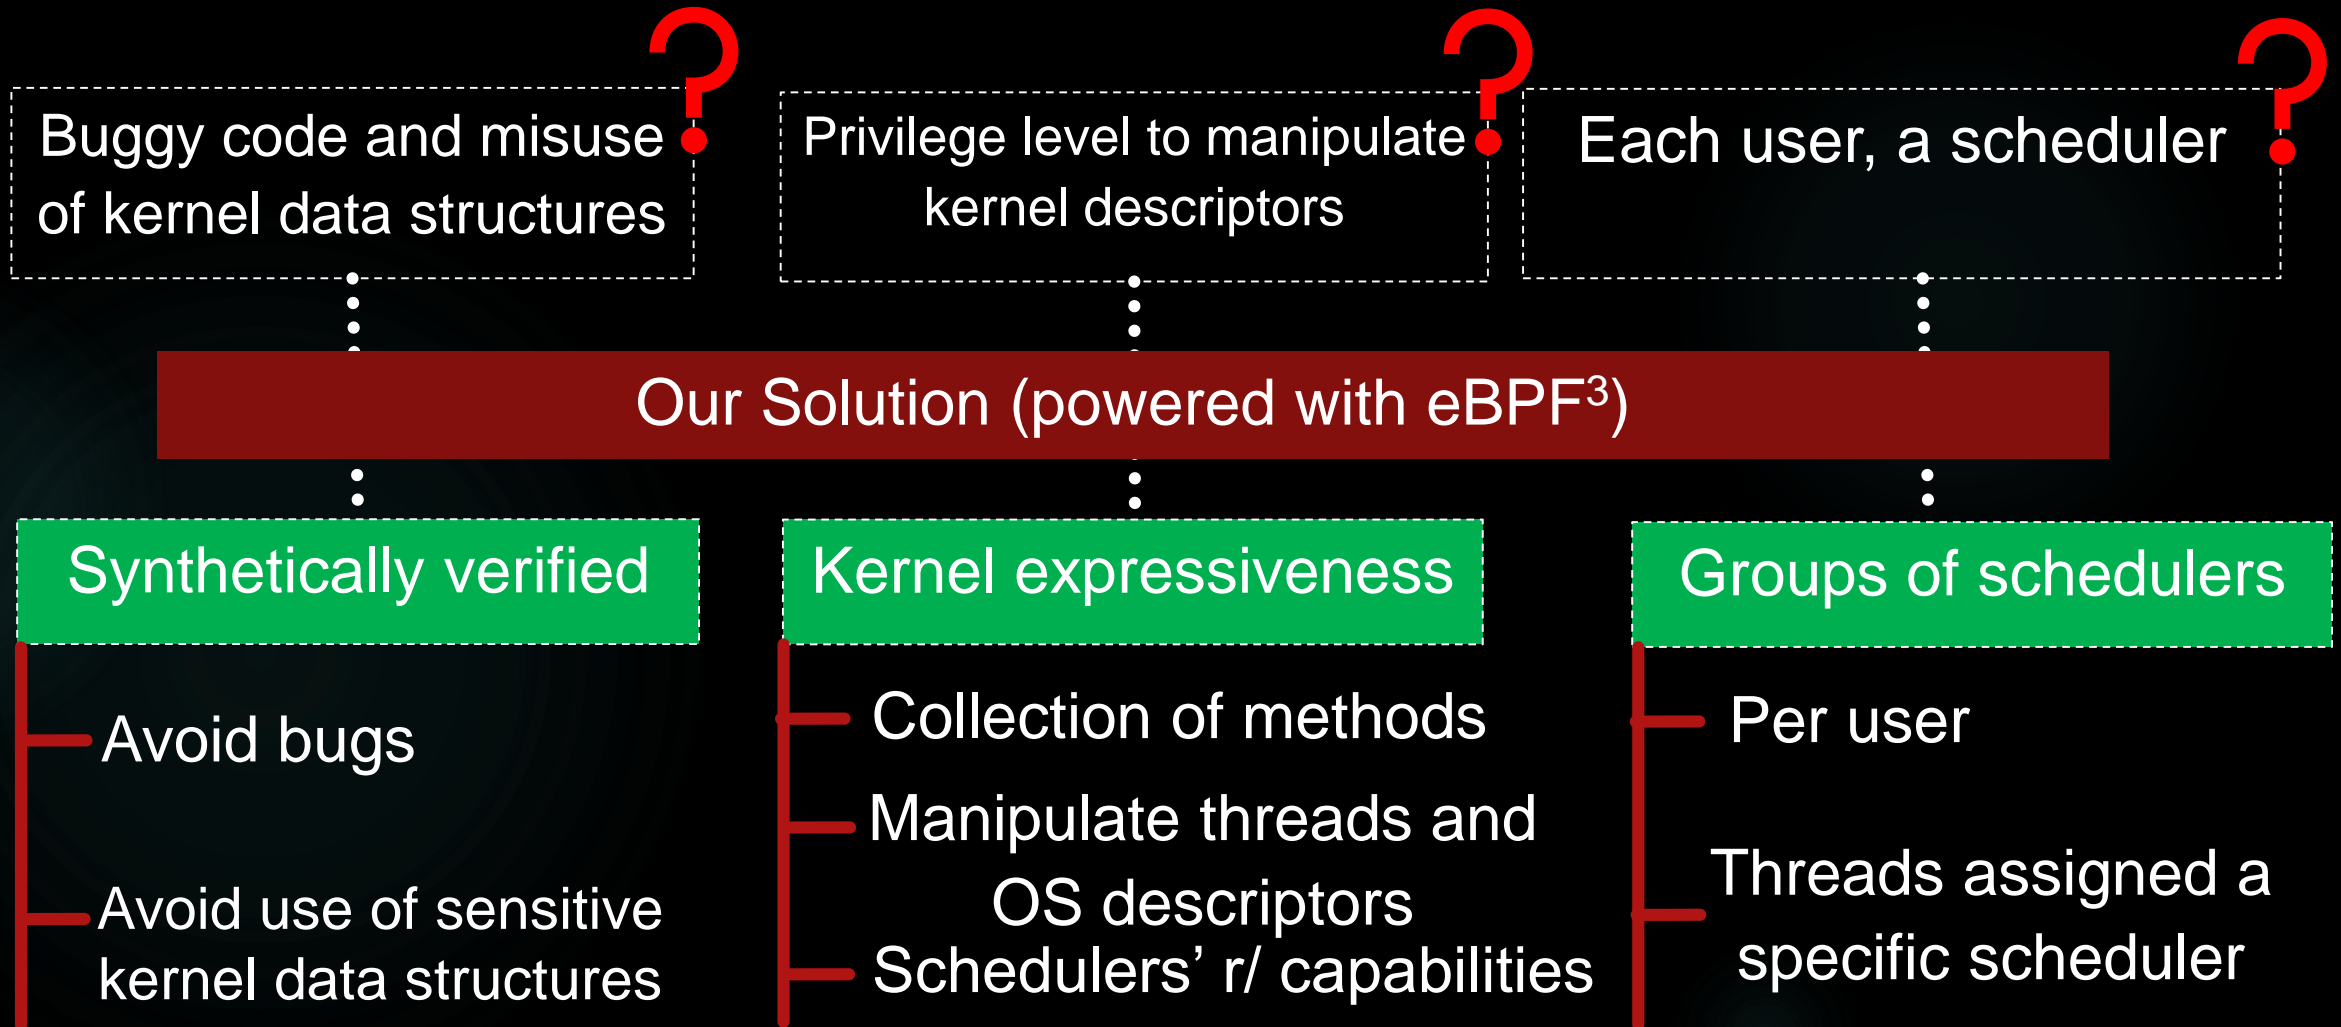
Buggy code and misuse  
of kernel data structures

Privilege level to manipulate  
kernel descriptors

Each user, a scheduler

# Our idea: Towards user-programmable schedulers

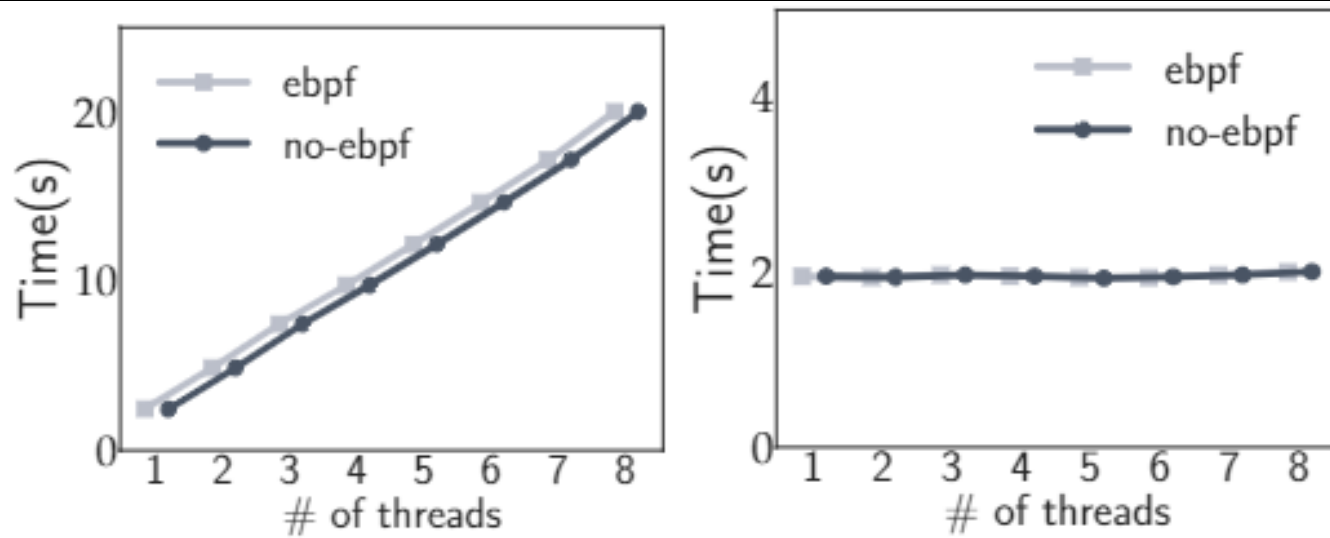
Users **express** the scheduler behavior at **runtime**





# Our idea: Towards user-programmable schedulers

Some **initial results** --- simple FIFO scheduler with our mechanism



(a) Execution time of a process.

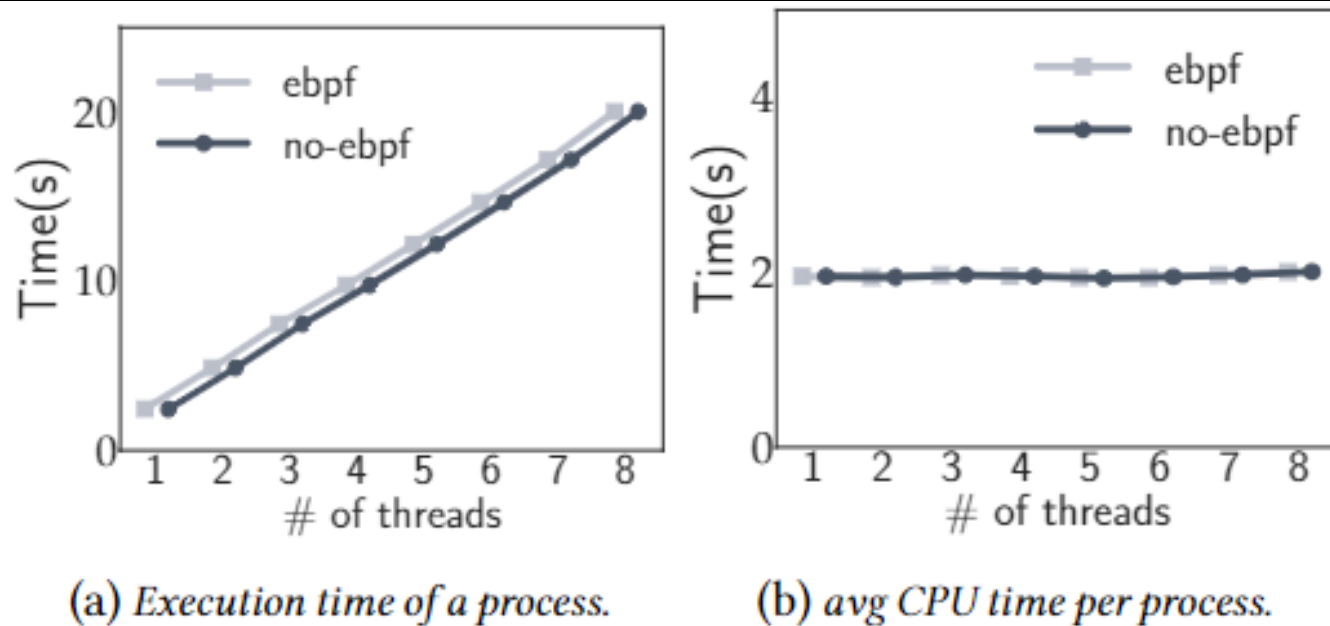
(b) avg CPU time per process.

**Figure 2.** eBPF-custom FIFO against Linux standard FIFO. We compute the (a) the execution time (wall-clock time) of one program, and (b) the average CPU time of each program as we increase the number of running programs.

Each thread computes  
prime numbers between  
**1** and **100,000**

# Our idea: Towards user-programmable schedulers

Some **initial results** --- simple FIFO scheduler with our mechanism

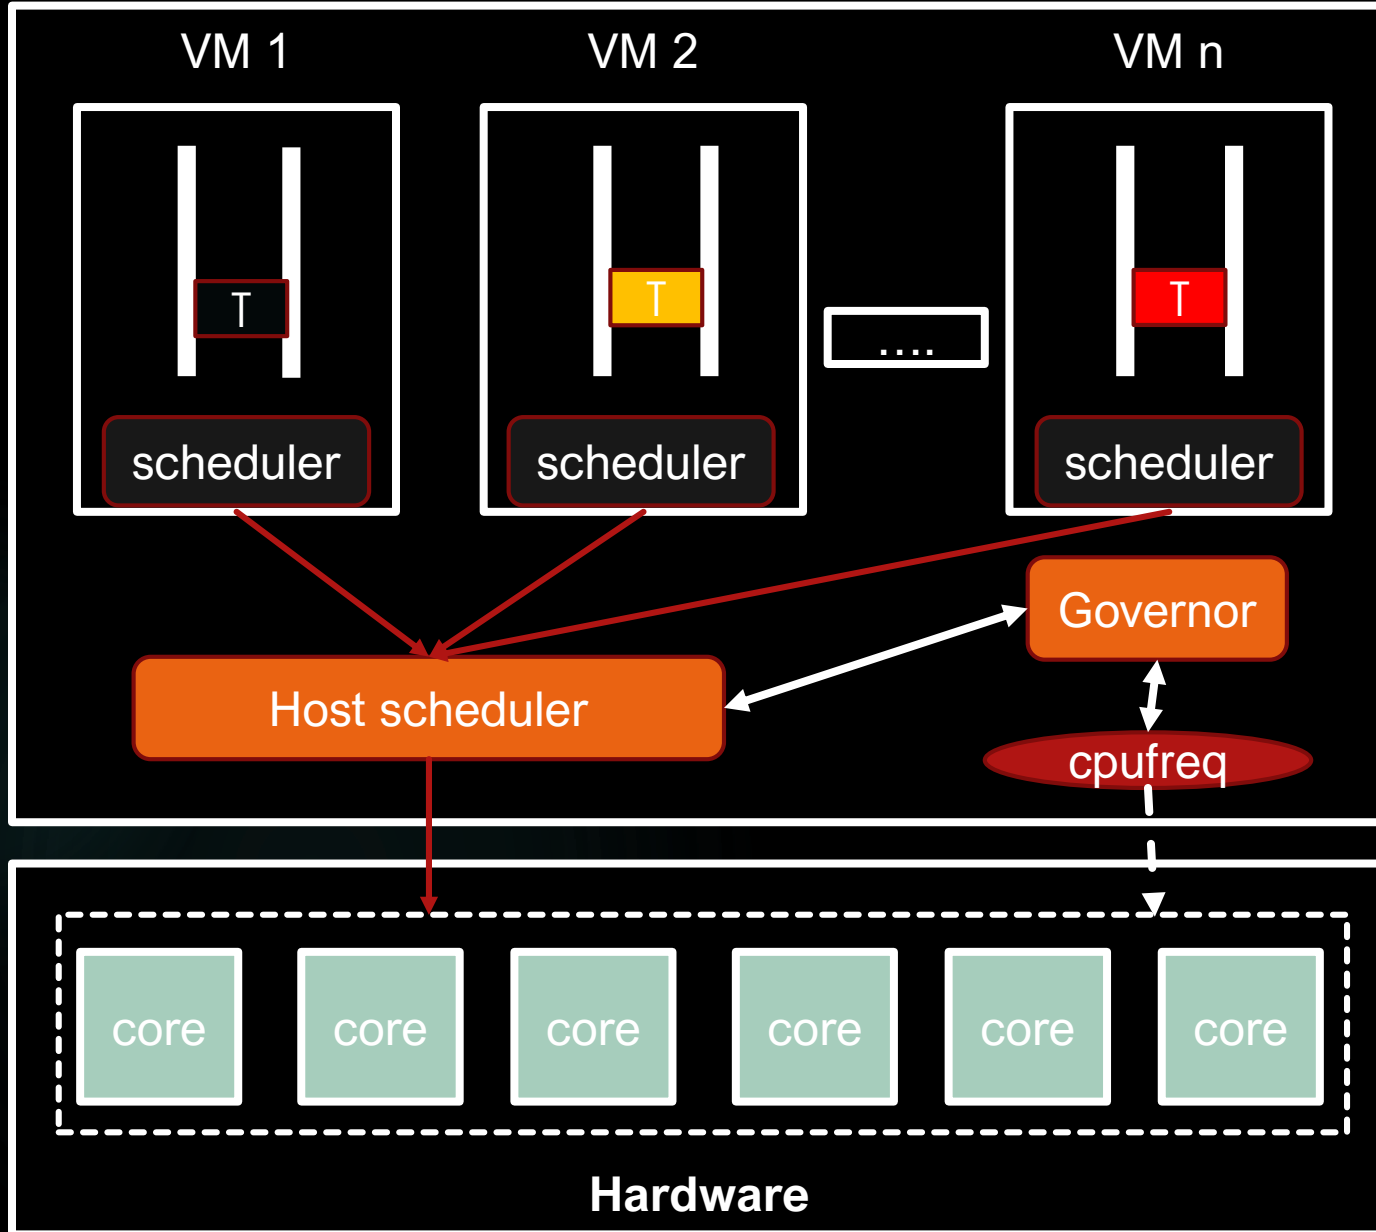


Interesting results but some work remains

- **JIT compiler performs poorly for loops**
- **Security issues regarding kernel data usage --- what to allow?**

**Figure 2.** *eBPF-custom FIFO against Linux standard FIFO. We compute the (a) the execution time (wall-clock time) of one program, and (b) the average CPU time of each program as we increase the number of running programs.*

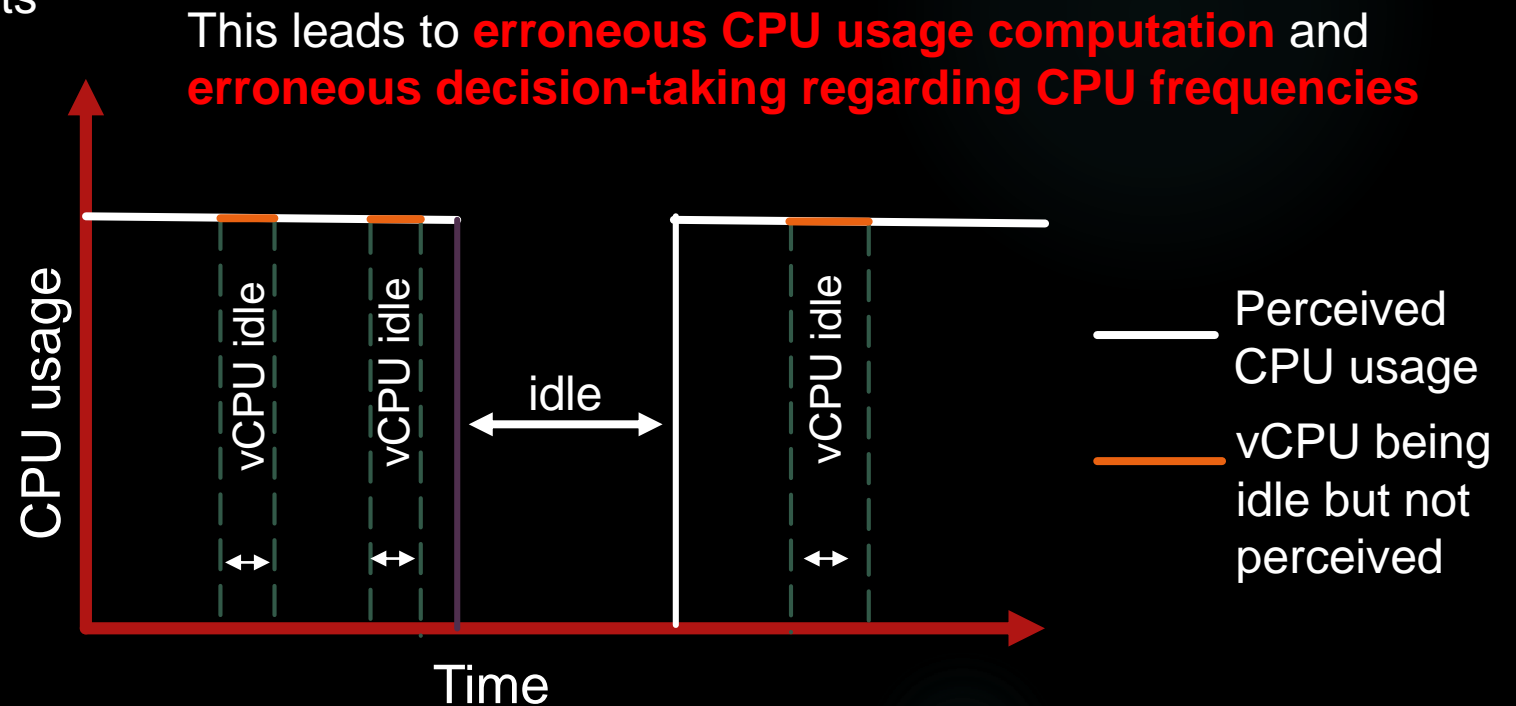
## PROBLEM 2: Incorrect accounting for IDLE ISOLATION UNITS



- ❑ Governors decide at which frequency the CPU should run favoring either performance or energy savings.
- ❑ Rely on metrics maintained by the scheduler to have an idea of each core CPU usage

## PROBLEM 2: Incorrect accounting for IDLE ISOLATION UNITS

- ❑ The host scheduler view of the vCPU is limited. Apart from explicit idle instructions (e.g., sleep), the vCPU is viewed as running.
- ❑ However, a running vCPU can be idle
  - ▶ Idle loop
  - ▶ Waiting for software interrupts
  - ▶ ...



## Example: On-demand algorithm for Xen

```
Curr_time = NOW()
```

```
Time_since_epoch=curr_time-prev_time
```

Get time and  
elapsed time since  
last call

For each cpu (j):

```
curr_idle_ns=get_cpu_idle_time(j)
```

```
idle_since_epoch=curr_idle_ns-j.prev_idle_ns
```

```
j.prev_idle_ns=curr_idle
```

Compute idle time  
since last call

```
if(time_since_epoch<idle_since_epoch) continue;
```

```
curr_freq = get_current_freq()
```

```
load = 100* (time_since_epoch - idle_since_epoch)/(time_since_epoch)
```

```
load_freq = load * curr_freq
```

Compute current load  
based on the idleness

```
if(load_freq > max_load_freq) max_load_freq = load_freq
```

Get the maximum percentage  
load across each cpu

```
if( load_freq > upper_threshold) push_to_max_frequency()
```

```
if( load_freq < lower_threshold) :
```

```
next_freq = load_freq/(threshold-10);
```

```
push_to_next_freq(next_freq);
```

Increase or decrease frequency --- when decreasing, tries to  
get the frequency that will not instantaneously trigger up policy

# Governors overview across hypervisors

## Xen

- Tracks for each vCPU the idle time
- Assumes the vCPU was running beside the idle time
- Basic threshold computation to decide the next frequency

## Linux/KVM

- For each run queue, tracks runnable and running time with a polynomial approximation based on the scheduling entity load (PELT)
- By that, they can leave out the idleness and focus on the running part
- Basic threshold computation to decide the next frequency

# Our idea - WIP

## Collaboration

- Add a tracker in each isolation unit that observes the main process and can identify idle periods
- Updates a flag via shared memory to tell the host scheduler that the corresponding isolation unit is idle and should be considered in the accounting.



# Our idea - WIP

## Collaboration mechanism

- Add a tracker in each isolation unit that observes the main process and can identify idle periods
- Updates a flag via shared memory to tell the host scheduler that the corresponding isolation unit is idle and should be considered in the accounting.

## Issues

- Identifying idle periods can require to peek into the VM
- Sharing introduces challenges regarding security --- Cloud platforms may not like
- What is the right behavior regarding shared memory? Event-based or periodic reading ...



**Exploring Hardware-based tracking --- extend HWP**



Thanks

Questions and ideas ? 